

DIGITIZATION OF HANDWRITTEN LEDGERS

AN INTEGRATED APPROACH TO IMPROVE HANDWRITTEN TEXT RECOGNITION

SUBMITTED IN PARTIAL FULFILLMENT FOR THE DEGREE OF MASTER OF SCIENCE

RODERICK MAJOOR
12852724

MASTER INFORMATION STUDIES
DATA SCIENCE
FACULTY OF SCIENCE
UNIVERSITY OF AMSTERDAM

SUBMITTED ON 23.05.2024

	UvA Supervisor
Title, Name	Dr. N.J.E. (Nanne) van Noord
Affiliation	MultiX Lab
Email	n.j.e.vannoord@uva.nl



1 ABSTRACT

2 Handwritten text recognition is a crucial step in digitizing historical
3 documents. In this work, we will explore the challenges associated
4 with a handwritten text recognition model using historical ledger
5 pages from the Bank of Amsterdam. We first conduct an error
6 analysis of the model, and based on that make suggestions on
7 where and how the performance of the model could be improved.
8 The main concern is the fact that the model is not able to capture the
9 layout of the pages. We perform post-processing steps to improve
10 the output of the model and we create a layout analysis method,
11 which aims to capture the distinctive tabular layout found in the
12 ledger pages. While earlier research mainly focuses on capturing
13 different elements on documents, such as text, tables and images,
14 our work specifically focuses on capturing table cells and contents
15 from a tabular structure. This is done by isolating the different
16 elements on the page using computer vision tools. **We find our**
17 **layout analysis method to be successful in capturing the different**
18 **table cells, but not always accurately matching the ground truth**
19 **bounding boxes.** Furthermore, we show that post-processing steps
20 improve the quality of the output of the model.

21 KEYWORDS

22 HTR, Layout analysis, Computer vision, Historical documents, Ob-
23 ject detection

24 GITHUB REPOSITORY

25 <https://github.com/roderickmajoer/Thesis>

26 1 INTRODUCTION

27 The Bank of Amsterdam was an early bank established in 1609, de-
28 scribed by some as the first central bank [20]. The bank played an
29 important role in the financial world of the 17th and 18th century.
30 Transactions made at the bank from 1650 to 1800 were recorded in
31 ledgers and are still preserved to this day. These ledgers contain
32 information about transactions, both debit and credit, of customers
33 of the bank. The ledgers can thus be very important in analyzing
34 the money flow at the time. Digitizing these ledgers may help in
35 preserving them in a much easier to analyze format which would
36 allow for more research into the contents of the ledgers. However,
37 these ledgers are all handwritten making it hard to retrieve the
38 information out of them on large scale. Current efforts to retrieve
39 the information are done by using Handwritten Text Recognition
40 (HTR) tools to extract the text in the ledgers. The problem is that
41 the currently used HTR techniques have too many inaccuracies to
42 be used effectively. Because of this, the handwritten text cannot
43 be recognized correctly and can thus not be digitized properly. To
44 improve digitization efforts, it is crucial to find out why some hand-
45 written text cannot be recognized, what kind of errors are made
46 during the HTR process and how the HTR process can be optimized
47 to improve these inaccuracies. The research question we wish to
48 answer is:

50 *How can the categorization of errors in handwritten ledger analysis
51 be used to enhance the identification and resolution of text recognition
52 inaccuracies?*

53 Subquestions belonging to this research question are:

- 55 • What factors contribute to text recognition errors in hand-
56 written ledgers?
- 57 • What are the different types of errors encountered in hand-
58 written text recognition?
- 59 • How can specific processing methods improve text detection
60 and segmentation in handwritten ledgers?
- 61 • What strategies can be implemented to mitigate common
62 text recognition errors in handwritten ledgers?

63 Since we perform a number of different experiments, our report
64 follows the following structure: we start by comparing some related
65 work. Then we will provide a description of the dataset. After that
66 we look at the different experiments, where for each we provide our
67 methodology including evaluation steps as well as results and dis-
68 cussion of that particular experiment. We end with a more general
69 discussion and conclusion.

70 2 RELATED WORK

71 The research gap being addressed in this project revolves around the
72 challenges associated with the digitization of handwritten ledgers.
73 The proposed research will contribute in closing the research gap
74 within the HTR domain, specifically looking at historical handwrit-
75 ten documents. The digitization of handwritten ledgers contains
76 several steps. In this section, we will look at key research papers
77 that discuss different approaches to the steps of our problem.

78 2.1 Document Layout Analysis

79 The biggest challenge in our digitization task is being able to find the
80 correct word order from the HTR output. Seuret describes layout
81 analysis and finding the correct word order as one of the main
82 challenges in HTR [21]. Several approaches to solve this problem
83 have been tried. These approaches are typically divided in either
84 deep learning methods or more classical computer vision methods.
85 We will show recent advancements made in both approaches.

86 Kastanas et al. show the use of different deep learning archi-
87 tectures for layout analysis [8]. Transformer-based, graph-based
88 models, and convolutional neural networks (CNN) are compared.
89 The CNN-based model YOLOv5 and transformed-based model Lay-
90 outLMv3 both show promising results for identifying most elements
91 on documents, with YOLOv5 performing slightly better [8].

92 Multiple other studies describe the use of transformer based
93 methods to perform layout analysis and optical character recog-
94 nition (OCR) [10, 13, 24]. Furthermore, Smock et al. show that
95 transformer-based object detection models can produce excellent
96 results for the tasks of detection, structure recognition, and functional
97 analysis of tables [22]. While these approaches show promising
98 results, they are not tested for our specific task at hand.

99 Oliveira et al. propose dhSegment, an open-source implementation
100 of a CNN-based pixel-wise predictor for document segmentation.
101 The approach aims to address various document processing
102 tasks simultaneously, including page extraction, baseline extraction,
103 layout analysis, and illustration and photograph extraction. They
104 show that a single CNN architecture can be used across tasks with
105 competitive results [2].

106 Drobny et al. propose a holistic method that applies Mask R-CNN
107 for text line extraction in historical documents. Their work achieved
108 state-of-the-art results on well known historical datasets [6]. While
109 their work does not directly transfer to our task, it does show the
110 potential for document segmentation and possibly table recognition
111 abilities of Mask R-CNN.

112 The aforementioned studies show that the use of deep learning
113 methods can be used to perform layout analysis as well as table
114 recognition tasks in historical documents. However, these methods
115 often require massive amounts of training data to be used effectively.
116 To use these methods, we should look whether pre-trained models
117 transfer well to our dataset and are able to detect the layout of our
118 documents. Other approaches use more classical computer vision
119 techniques to segment a document in order to recognize the layout.
120 These methods often require homogeneity of the documents to
121 work effectively.

122 Lehenmeier et al. show a method that combines various state-
123 of-the-art methods for OCR pro- cessing to detect layout on a
124 document [9]. To perform table recognition, they make use of tech-
125 niques such as binarization, denoising, and Hough line detection
126 to find relevant parts of the tables. By taking the intersection of
127 relevant horizontal and vertical Hough lines they are able to find
128 table cells and thus determine the layout of the table [9].

129 Bulacu et al. introduce a method based on contour tracing that
130 generates curvilinear separation paths between text lines in order to
131 preserve the ascenders and descenders of text lines to determine the
132 layout in handwritten historical documents [3]. With this approach
133 they are able to determine seperate elements of tables in historical
134 documents.

135 Liu et al. describe the use of several key steps, including skew
136 correction, region segmentation, and page layout analysis on his-
137 torical Tibetan documents [12]. Skew correction is achieved by
138 leveraging baseline features and Hough transform to determine
139 the document's skew angle. Subsequently, region segmentation is
140 accomplished by identifying borders through a series of preprocess-
141 ing steps, including median filtering, Gaussian smoothing, Sobel
142 edge detection, and removal of small area regions. These processes
143 collectively enable accurate positioning of document borders for
144 further analysis and structure extraction [12].

145 Liang et al. describe the use of Hessian and Gabor filters in
146 order to find table structures on a document [11]. They show the
147 possibility for column segmentation for the tables using the found
148 lines.

149 Prieto et al. address the challenge of document image understand-
150 ing in documents with complex layouts like tables. They compare
151 two approaches and show promising results [19]. Their approaches
152 take the text lines outputted by HTR systems and use machine

153 learning methods to group these text lines and classify cells based
154 on these grouped text lines in order to find the tabular structure.

155 It can be seen that depending on the dataset and the homogene-
156 ity of the documents, different approaches for layout analysis/table
157 recognition may be suitable. When a similar table structure is clearly
158 drawn on all the documents, approaches using basic image process-
159 ing techniques could be suitable to determine the structure of the
160 table. When the structure is not easily identifiable on the page, it
161 might be better to look at approaches that use the output of HTR
162 systems such as text lines or words locations to classify table cells.
163 Machine learning based approaches could also be useful in these
164 cases but often require many training data.

165 2.2 HTR Processing Techniques

166 Typical HTR pipelines consist of pre-processing an image to make
167 it ready for HTR, running the image through the HTR system and
168 post-processing the output of the HTR system [9].

169 Studies have shown that pre-processing images can lead to better
170 results when using them as input for document layout analysis or
171 OCR. Thus, we want to optimize our image material such that the
172 document layout analysis and HTR can be performed more accu-
173 rately. There are different methods for doing this such as adjusting
174 the contrast to remove artifacts and noise or fixing the image align-
175 ment [1, 5, 9]. It is crucial to find out which pre-processing steps
176 might be useful to our problem at hand. Chen et al. show that the
177 use of several pre-processing steps including character segmenta-
178 tion, tilt correction, offset correction, size normalization and image
179 thinning lead to better HTR output results [4].

180 Post-processing OCR and HTR output typically consists of trying
181 to correct misidentified words or characters. One strategy to do
182 this is to create a dictionary of candidate characters and use the
183 Levenshtein distance to calculate the distance between predicted
184 and dictionary characters [18]. The use of a dictionary for allowed
185 characters is shown to be a good strategy to eliminate HTR errors
186 [9]. Other studies show the use of language models to correct word
187 output and spelling in HTR systems [15, 18].

188 2.3 HTR Evaluation Strategies

189 Typically, HTR systems are evaluated using the word error rate
190 (WER) and character error rate (CER). However, the evaluation
191 of page-level HTR output faces challenges primarily due to the
192 Reading Order (RO) problem [23]. In their study, they define a bag-
193 of-words Word Error Rate (bWER) to accurately measure page-level
194 RO-independent word errors as well as a regularized version of the
195 Hungarian Algorithm (HA) to compute word- and character-level
196 RO-independent recognition accuracy [23]. Another strategy is
197 the use of the intersection over union (IoU) score which matches
198 ground truth and HTR output boxes based on their coordinates
199 on the page [17]. It is then possible to use the WER and CER to
200 calculate the accuracy.

201 3 DATA DESCRIPTION

202 The dataset used in this study consists of pages from the collection
203 of ledgers from the Bank of Amsterdam. The full collection can be
204 found on the [website](#) of The Amsterdam City Archives. In the series

of ledgers, the current accounts of the customers were maintained. Each customer had one or more pages, with smaller traders having a portion of a page. An alphabetical list of the traders and the page of their current account can be found in the index. The ledgers kept track of the amounts that were debited and credited from and to a customer. The ledger pages consist of several columns containing information such as the date, account holder name, account number and amount debited/credited. Figure 9 shows an example of what a page looks like. Our dataset is split into multiple subsets that come from different books and that all have slightly varying pages, while still having structural similarity. Figure 1 shows the amount of pages for each subset in the dataset. The total dataset consists of 68 pages. A word count for each page is shown in Figure 2

The ground truth and HTR system output both use a PageXML format to store values on the page. These PageXML store the coordinates for each item (e.g. textregion, tableregion, etc.) on the page. We can plot the values in the PageXML file over our image to better show this. Figure 10 shows the PageXML data of the ground truth and HTR system plotted on the original image for part of one page in our dataset.





Figure 1: The amount of pages for each subset in our dataset.



$$CER = \frac{S_c + D_c + I_c}{N_c} \quad (2)$$

Where S_c , D_c , I_c are the number of character substitutions, deletions and insertions respectively and N_c is the total number of characters in the ground truth.

We will also calculate recall and precision scores. To calculate those for HTR output, we first need to define what constitutes true positives (TP), false positives (FP), and false negatives (FN) in our context.

- **TP (True Positive):** Characters correctly recognized by the HTR system.
- **FP (False Positive):** Characters incorrectly recognized by the HTR system (i.e., system identifies text where there is none or incorrectly identifies text).
- **FN (False Negative):** Characters that are present in the ground truth but not recognized by the HTR system.

The recall and precision are then:

$$Recall = \frac{TP}{TP + FN} \quad (3)$$

$$Precision = \frac{TP}{TP + FP} \quad (4)$$

4.3 Results

The accuracy of the current HTR system using the aforementioned metrics is shown in Table 1. As can be seen, the accuracy of the base model is not flawless. The recall is quite high, showing that not a lot of false negatives are found. However, the precision is lower, showing that the models incorrectly captures text. A more in-depth analysis about the amount of character insertions, deletions, substitutions and matches made by the HTR system can be seen in Figure 4 and Figure 3. Here, we can clearly see that substitutions and mainly insertions of characters are the main problem of the HTR model while deletions are relatively low. In order to improve the model, it is thus crucial to find a way to mainly reduce the insertions and substitutions made by the model. Figure 5 shows the frequencies of characters inserted, substituted and deleted. Here we can see that a lot of 'invalid' characters are inserted. These are non-alphanumeric characters that should not be part of our dataset. Furthermore, we see that many numbers and letters that look alike (such as i and 1) are substituted. We use these analyses and the fact that the HTR system is not able to capture the layout of the pages to conduct a list of errors made by the HTR system. A selection of the most common errors and possible solutions is shown in Table 2. In further sections, we will use the results of this section to improve the model.

Table 1: Baseline scores for the currently used HTR system. We matched the ground truth words with HTR output words based on highest IoU value.

Metric	Value
WER	0.4243
CER	0.3422
Recall	0.9475
Precision	0.7403
Total Edit Distance	18920

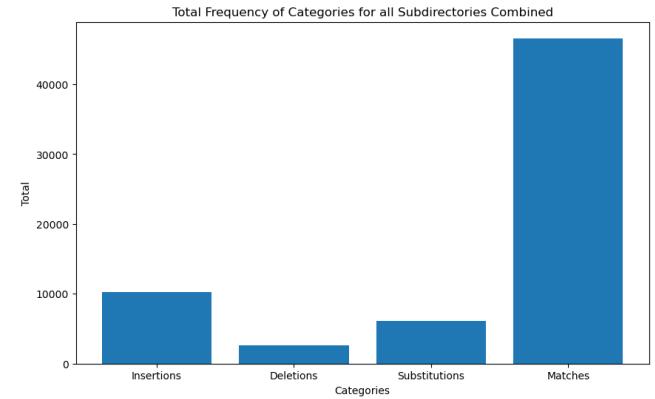
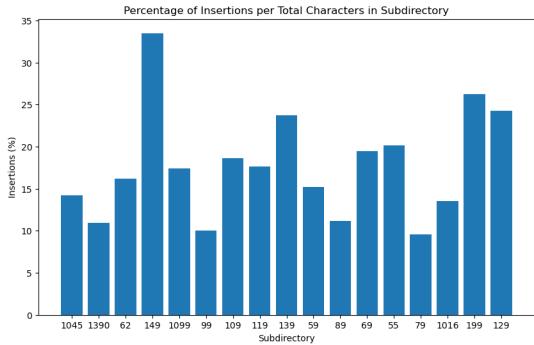
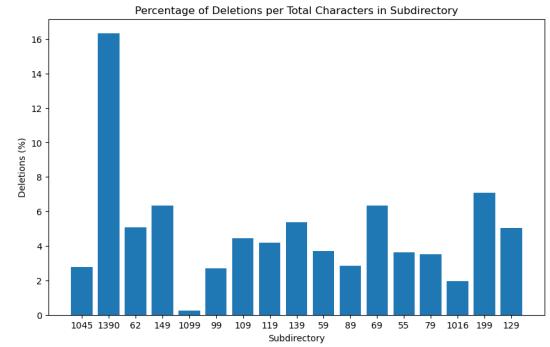


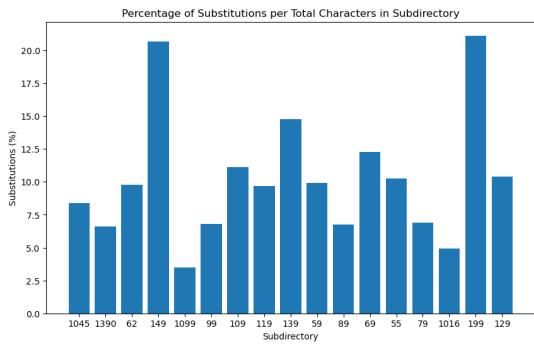
Figure 3: The total amount of insertions, deletions, substitutions and matches made by the HTR system compared to the ground truth values. The values are on a character level and made word per word.



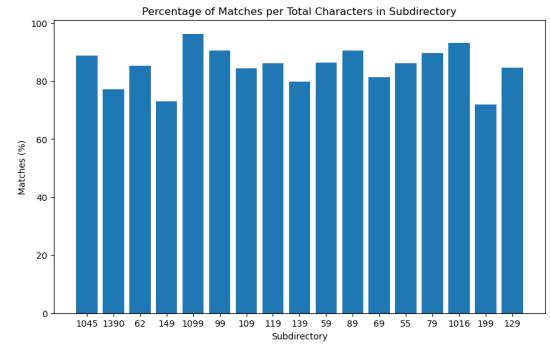
(a) Amount of insertions made by the HTR system per subdirectory.



(b) Amount of deletions made by the HTR system per subdirectory.



(c) Amount of substitutions made by the HTR system per subdirectory.




(d) Amount of matches made by the HTR system per subdirectory.

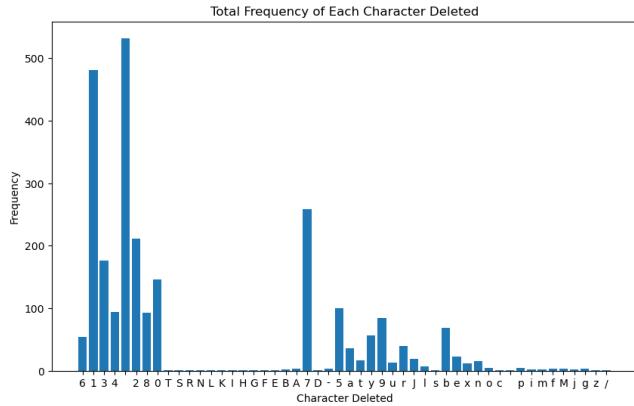
Figure 4: The amount of insertions, deletions, substitutions and matches made by the HTR system compared to the ground truth values. The values are on a character level and made word per word and as a percentage of total characters in ground truth.

Table 2: Categorization of HTR System Errors

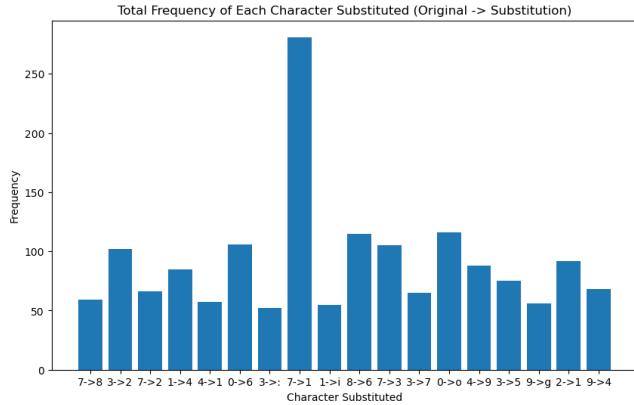
Error Type	Description	Possible Solution
Deletions	Characters not recognized at all (false negative)	Improve HTR system, post-processing
Insertions	Characters recognized that are not there (false positives)	Improve HTR system, post-processing
Substitutions	Characters recognized as another character	Improve HTR system, post-processing
Layout - Word Split	Single word is split into separate words	Layout analysis
Layout - Word Merge	Multiple words are recognized as one word	Layout analysis
Layout - Data Split	Data wrongly split (e.g., 'feb 5' in GT becomes 'feb' and '5' in HTR)	Layout analysis
Layout - Word Order	Word order not recognized	Layout analysis
HTR - Bounding Box	Wrong location bounding box vs actual location word	Improve HTR system
Case Sensitivity	Case sensitivity issues	Post-processing
Invalid Characters	Characters that cannot occur	Post-processing



(a) Frequency of inserted characters made by the HTR system.



(b) Frequency of deleted characters made by the HTR system.



(c) Frequency of substituted characters made by the HTR system. The substitutions are shown as (original character -> substituted character). We only show the substitutions with a frequency of 50 or higher.

Figure 5: The frequencies of inserted, deleted and substituted characters made by the HTR system.

5 POST-PROCESSING

5.1 Method

From our error analysis, we find that using post-processing techniques may lead to improved accuracy. Post-processing in our experimental framework focuses on refining the output generated by the HTR system. One such post-processing step involves the removal of characters from the output strings that do not align with the ground truth dictionary [9, 18]. Analysis of the HTR output (as depicted in Figure 5a) reveals several insertions that are not valid letters or numeric characters. To address this, we employ regular expressions (regex) to filter out invalid characters from the HTR output. We are then left with alphanumeric characters only. Furthermore, the HTR system often substitutes letters for numbers, as can be seen in Figure 5c. We can undo these substitutions in some cases. For instance, when a single character is found to be a letter by the HTR system, we can expect this to be wrong, since single character words are never letters in the ground truth. Also, numbers that contain a letter (such as '11i0') are probably wrong. We use these facts and the most common made substitutions as found in our earlier analysis to recover the right characters. We use the same evaluation metrics as before to compare our method to the baseline scores.

5.2 Results

We compare the post-processed HTR output to the original HTR output to see the different accuracy scores. The accuracy scores for both the baseline and post-processed HTR output are shown in Table 3. The comparison of the performances is shown in Figure 6. We see a very slight decrease in recall while all other metrics improve. This shows that our post-processing strategy is successful in eliminating some of the false positives (insertions and substitutions) of the original model while maintaining almost the same rate of false negatives (deletions). However, we can still see that the precision is relatively low compared to the recall, showing that there is still a good amount of false positives that our method did not catch.

Table 3: WER, CER, precision and recall scores for the HTR system output with and without post-processing techniques applied.

Metric	Baseline	Post-Process
WER	0.4243	0.3916
CER	0.3422	0.3169
Recall	0.9475	0.9418
Precision	0.7403	0.7622
Total Edit Distance	18920	17508

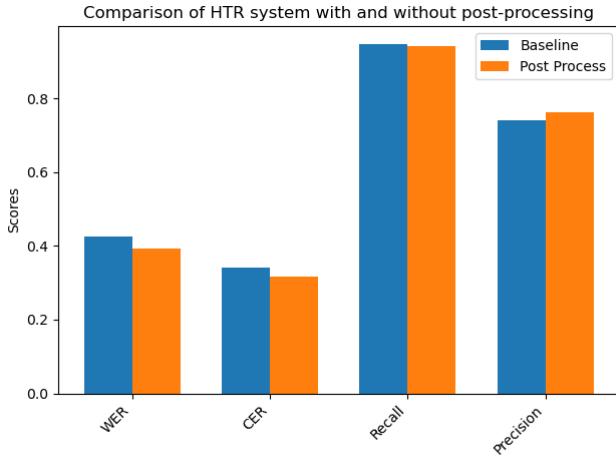


Figure 6: The WER, CER, precision, and recall scores for the HTR system with and without post-processing applied.

We then apply the Hough Line Transform to detect the lines in the image. This can be used to detect (nearly) straight lines in images [9]. We filter the detected lines based on their length and position. For each detected line, we calculate its slope and intercept, and then draw an extended line from the top to the bottom of the foreground area (which contains the ledger page) in the image. This is done using basic principles of line equations in coordinate geometry.

Finally, we perform additional dilation and erosion to merge close lines and remove final noise. We then find contours in the image and draw a line from the top to the bottom of each contour. We are then left with an image containing the contours of each column section, which can then be used to create bounding boxes for the different columns. The whole process is shown in Figure 12.

To detect the rows present in the ledger pages, we have to use a different strategy, since there are no physical markings for row lines on the pages. We will use the word bounding boxes which are outputted by the HTR system in order to find the row bounding boxes.

We start by extracting the bounding boxes of the words outputted by the HTR system. We will only use the word bounding boxes that are kept after performing the post-processing. Since the images in our dataset typically consist of two tables (debit and credit) we have to determine where both tables start and end. We do this by finding the middle column line from the list of column lines we found by our method above. We then separate the bounding boxes into two groups: left boxes and right boxes, based on their x-coordinates relative to the middle column line.

We then take a list of boxes (either left boxes or right boxes) and group them into rows. This is done by sorting the boxes by their x-coordinates (from left to right) and then grouping boxes with similar y-coordinates together. For each box, we find the best row to append the box to, based on the minimum difference in y-coordinates. If no suitable row is found, we create a new row. The result of these steps are shown in Figure 13.

Once we have found the rows, we can extend the bounding boxes from the left side of the ledger page to the middle for the left boxes and from the middle to the right side of the page for the right boxes. We are now left with both the column and row bounding boxes.

Once the row and column bounding boxes are determined, we can use the intersection of these bounding boxes to determine the table cells. Figure 7 shows an example of these found table cells.

6.2 Evaluation

For the layout analysis methods, we take the ground truth table cell bounding boxes and compare them to our method's predicted table cell bounding boxes. We do this by matching each ground truth cell to a predicted cell based on the highest IoU score between the two. Only one ground truth cell can be matched to one predicted cell and vice versa. We then determine whether cells are a true positive by setting an IoU threshold. If ground truth and predicted cells are matched with an IoU score higher than this threshold, we consider it a true positive. Ground truth cells that are not matched with an IoU score above the threshold are considered a false negative and

6 LAYOUT ANALYSIS

6.1 Method

Up until this point, we have been comparing the HTR output to the ground truth values by matching their bounding boxes based on the highest IoU score. We did this because the HTR system currently is not able to capture the layout of the tabular structure of the ledger pages and thus we could not just match the output directly, since the word order would be wrong. The final part of the methodology involves analyzing the layout of the documents. Recognizing the layout is a crucial step in digitizing the ledger pages, since without it there is no good way to use the output of the HTR system because the correct word order would not be known. We will perform the layout analysis by using image processing techniques as well as making use of the HTR output. For this, we will use the [OpenCV \(cv2\)](#) library to perform the necessary processing steps.

To detect the tabular structure on the ledger pages, we must find a way to detect the columns and rows of the table to get the table cells. The ledger pages contain physical lines, which separate the different columns. We start by trying to detect these lines.

The first step is to pre-process our image. We do this by converting our image to grayscale and applying a Gaussian blur on the grayscale image. This helps reducing the computational complexity and reduces high frequency noise [14].

Next, we apply two types of thresholding to the blurred image: Otsu's thresholding and Adaptive Gaussian thresholding. Otsu's thresholding only captures the most clear markings on the page, such as the text, while the more thin and faint column lines are not captured [3]. The Adaptive Gaussian threshold captures both the text as well as the column lines. By subtracting the first threshold from the latter, we are left with the column lines and some noise.

Following thresholding, we perform morphological operations, specifically dilation and erosion [16]. We use specific kernels to remove noise while retaining the vertical lines.

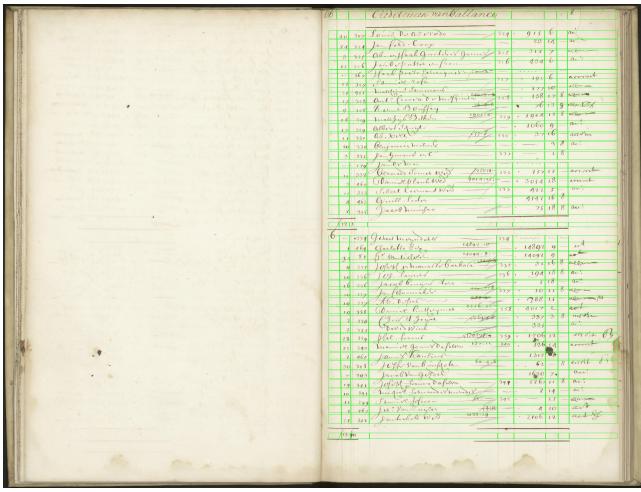


Figure 7: The bounding boxes of the predicted table cells after taking the intersection of the rows and columns.

matched with predicted cells while this is not the case when using a higher IoU threshold. One of the main reasons for this behaviour is the fact that the bounding boxes of our predicted cells are much more tight than the ground truth boxes. Furthermore, the bounding boxes for our predicted table cells are often placed slightly lower than the ground truth bounding boxes. This is because we use the word bounding boxes from the HTR model to determine the rows, and these bounding boxes are almost always found to be slightly lower than the actual location of the words on the image. This causes the IoU score of a predicted cell and ground truth cell to be relatively low, even when both capture the same contents of the page. When taking a IoU threshold of 0.25, as shown in Table 7, we see the best ratio of IoU scores and accuracy scores. Here we see that matched cells have a mean IoU score of 0.473, which is quite close to the usual threshold of 0.5. Furthermore, we see a score of 0.804, 0.898 and 0.844 for mean precision, recall, and F1-score respectively. This shows that we have a good rate of false positives and false negatives, as well as having a reasonable area overlap for predicted and ground truth cells.

The accuracy of comparison of text contents for ground truth and predicted cells can be seen in Table 5. Here, 'LT 0' until 'LT 0.75' means how much percent of the text in the ground truth and predicted cells differ. So for the first row, 0.5369 means that 53.69% of texts are matched exactly, 0.5996 means that 59.96% of texts are matched with at most 25% of characters changed etc. If we look at the IoU threshold of 0.25, which seemed to perform well in the previous result, we can see that about 85% of words are matched with at most 75% of the characters differing. Furthermore we see a CER of 0.3859. This is higher than the CER found in Section 4. However, this can partly be explained due to the fact that in the first experiment, we calculated CER by matching one HTR word with a ground truth cell text. That means that we discarded some of the false positives (HTR words that were not matched). In this case, we first append all HTR words into a predicted cell and then compare predicted cell content to ground truth cell content, meaning that in this case we do keep these false positives when comparing the texts. Keeping this in my mind, we can determine that most of the predicted cells and ground truth that are matched by IoU, are also matched based on text contents, showing the success of our method.

Table 4: Mean table cell performance metrics at different IoU Thresholds. 'IoU Score' indicates the mean score for the true positives found at this IoU threshold.

IoU Threshold	IoU Score	Precision	Recall	F1
>0	0.182	0.845	0.946	0.887
0.25	0.473	0.804	0.898	0.844
0.5	0.627	0.416	0.459	0.434

63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509




Figure 8: The mean IoU, precision, recall and F1 scores for ground truth and predicted table cell bounding boxes for different IoU thresholds.

Table 5: Mean accuracy and CER for matching ground truth table cells with predicted table cells using different IoU thresholds and comparing their text contents using normalized Levenshtein distance thresholds. Values for LT 0, 0.25, 0.50, and 0.75 indicate the mean accuracy at which the normalized Levenshtein distance between the ground truth and predicted texts is within 0 (total match), 25%, 50%, and 75% respectively. CER indicates the average character error rate for the text content in the matched cells.

IoU Thresh.	LT 0	LT 0.25	LT 0.50	LT 0.75	CER
>0	0.5369	0.5996	0.7333	0.8242	0.4290
0.25	0.5555	0.6213	0.7563	0.8469	0.3859
0.50	0.5898	0.6586	0.7984	0.8770	0.3362

7 DISCUSSION

Write your discussion here. Do not forget to use sub-sections. Normally, the discussion starts with comparing your results to other studies as precisely as possible. The limitations should be reflected upon in terms such as reproducibility, scalability, generalizability, reliability and validity. It is also important to mention ethical concerns.

8 CONCLUSION

Write your conclusion here. Be sure that the relation between the research gap and your contribution is clear. Be honest about how limitations in the study qualify the answer on the research question.

REFERENCES

- [1] Yasser Alginahi. 2010. *Preprocessing Techniques in Character Recognition*. <https://doi.org/10.5772/9776>
- [2] Sofia Ares Oliveira, Benoit Seguin, and Frederic Kaplan. 2018. dhSegment: A Generic Deep-Learning Approach for Document Segmentation. In *2018 16th International Conference on Frontiers in Handwriting Recognition (ICFHR)*. IEEE. <https://doi.org/10.1109/icfhr-2018.2018.00011>
- [3] Marius Bulacu, Rutger van Koert, Lambert Schomaker, and Tijn van der Zant. 2007. Layout Analysis of Handwritten Historical Documents for Searching the Archive of the Cabinet of the Dutch Queen. In *Ninth International Conference on Document Analysis and Recognition (ICDAR 2007)*, Vol. 1. 357–361. <https://doi.org/10.1109/ICDAR.2007.4378732>
- [4] Tieming Chen, Guangyuan Fu, Hongqiao Wang, and Yuan Li. 2020. Research on Influence of Image Preprocessing on Handwritten Number Recognition Accuracy. In *The 8th International Conference on Computer Engineering and Networks (CENet2018)*, Qi Liu, Mustafa Misir, Xin Wang, and Weiping Liu (Eds.). Springer International Publishing, Cham, 253–260.
- [5] Sergio Correia and Stephan Luck. 2023. Digitizing historical balance sheet data: A practitioner’s guide. *Explorations in Economic History* 87 (2023), 101475. <https://doi.org/10.1016/j.eeh.2022.101475> Methodological Advances in the Extraction and Analysis of Historical Data.
- [6] Ahmad Drobly, Berat Kurar Barakat, Reem Alaasam, Boraq Madi, Irina Rabaev, and Jihad El-Sama. 2022. Text Line Extraction in Historical Documents Using Mask R-CNN. *Signals* 3, 3 (2022), 535–549. <https://doi.org/10.3390/signals3030032>
- [7] Wafaa S. El-Kassas, Cherif R. Salama, Ahmed A. Rafea, and Hoda K. Mohamed. 2021. Automatic text summarization: A comprehensive survey. *Expert Systems with Applications* 165 (2021), 113679. <https://doi.org/10.1016/j.eswa.2020.113679>
- [8] Sotirios Kastanas, Shaomu He, and Yi He. 2023. Document AI: A Comparative Study of Transformer-Based, Graph-Based Models, and Convolutional Neural Networks For Document Layout Analysis. arXiv:2308.15517 [cs.CL]
- [9] Constantin Lehnenmeier, Manuel Burghardt, and Bernadette Mischka. 2020. *Layout Detection and Table Recognition – Recent Challenges in Digitizing Historical Documents and Handwritten Tabular Data*. Springer International Publishing, 229–242. https://doi.org/10.1007/978-3-030-54956-5_17
- [10] Minghao Li, Tengchao Lv, Jingye Chen, Lei Cui, Yijuan Lu, Dinei Florencio, Cha Zhang, Zhoujun Li, and Furu Wei. 2022. TrOCR: Transformer-based Optical Character Recognition with Pre-trained Models. arXiv:2109.10282 [cs.CL]
- [11] Xusheng Liang, Abbas Cheddad, and Johan Hall. 2021. Comparative Study of Layout Analysis of Tabulated Historical Documents. *Big Data Research* 24 (2021), 100195. <https://doi.org/10.1016/j.bdr.2021.100195>
- [12] Huaming Liu, Xuehui Bi, and Weilan Wang. 2019. Layout analysis of historical Tibetan documents. In *2019 2nd International Conference on Artificial Intelligence and Big Data (ICAIBD)*, 74–78. <https://doi.org/10.1109/ICAIBD.2019.8837040>
- [13] Tengchao Lv, Yuyan Huang, Jingye Chen, Lei Cui, Shuming Ma, Yaoyao Chang, Shaohan Huang, Wenhui Wang, Li Dong, Weiyao Luo, Shaoxiang Wu, Guoxin Wang, Cha Zhang, and Furu Wei. 2023. Kosmos-2.5: A Multimodal Literate Model. arXiv:2309.11419 [cs.CL]
- [14] Siddharth Misra and Yaokun Wu. 2020. Chapter 10 - Machine learning assisted segmentation of scanning electron microscopy images of organic-rich shales with feature extraction and feature ranking. In *Machine Learning for Subsurface Characterization*, Siddharth Misra, Hao Li, and Jiabo He (Eds.). Gulf Professional Publishing, 289–314. <https://doi.org/10.1016/B978-0-12-817736-5.00010-7>
- [15] Arthur Flor de Sousa Neto, Byron Leite Dantas Bezerra, and Alejandro Héctor Toselli. 2020. Towards the Natural Language Processing as Spelling Correction for Offline Handwritten Text Recognition Systems. *Applied Sciences* 10, 21 (2020). <https://doi.org/10.3390/app10217711>
- [16] OpenCV. 2024. Morphological Transformations. https://docs.opencv.org/3.4/dd/d7/tutorial_morph_lines_detection.html Accessed: 2024-04-18.
- [17] Rafael Padilla, Sergio Netto, and Eduardo da Silva. 2020. A Survey on Performance Metrics for Object-Detection Algorithms. <https://doi.org/10.1109/IWSSIP48289.2020>
- [18] Rémi Petitpierre, Marion Kramer, and Lucas Rappo. 2023. An end-to-end pipeline for historical censuses processing. *International Journal on Document Analysis and Recognition (IJDAR)* 26, 4 (March 2023), 419–432. <https://doi.org/10.1007/s10032-023-00428-9>
- [19] Jose Ramón Prieto, José Andrés, Emilio Granell, Joan Andreu Sánchez, and Enrique Vidal. 2023. Information extraction in handwritten historical logbooks. *Pattern Recognition Letters* 172 (2023), 128–136. <https://doi.org/10.1016/j.patrec.2023.06.008>
- [20] Stephen Quinn and William Roberds. 2009. *An economic explanation of the early Bank of Amsterdam, debasement, bills of exchange and the emergence of the first central bank*. Cambridge University Press, 32–70.
- [21] Mathias Seuret. [n.d.]. *Layout Analysis in Handwritten Historical Documents*. Chapter Chapter 4, 45–65. https://doi.org/10.1142/9789811203244_0004 arXiv:https://www.worldscientific.com/doi/pdf/10.1142/9789811203244_0004
- [22] Brandon Smock, Rohith Pesala, and Robin Abraham. 2022. PubTables-1M: Towards Comprehensive Table Extraction From Unstructured Documents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 4634–4642.
- [23] Enrique Vidal, Alejandro H. Toselli, Antonio Ríos-Vila, and Jorge Calvo-Zaragoza. 2023. End-to-End page-Level assessment of handwritten text recognition. *Pattern Recognition* 142 (2023), 109695. <https://doi.org/10.1016/j.patcog.2023.109695>
- [24] Yiheng Xu, Tengchao Lv, Lei Cui, Guoxin Wang, Yijuan Lu, Dinei Florencio, Cha Zhang, and Furu Wei. 2021. LayoutXML: Multimodal Pre-training for Multilingual Visually-rich Document Understanding. arXiv:2104.08836 [cs.CL]

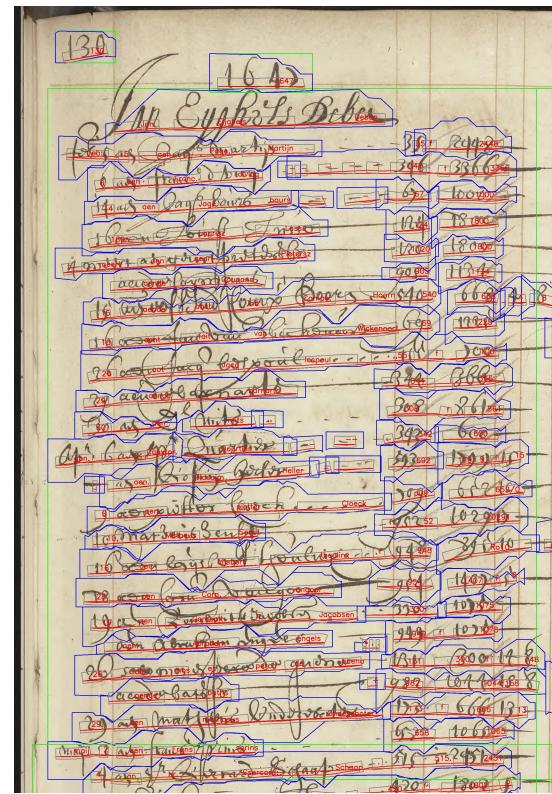
Figure 9: An example page from the ledger collection.

130 164

Mr. Euphile's Barber

1. Mr. Euphile's Barber	310	2442
2. Mr. Euphile's Barber	2442	35663
3. Mr. Euphile's Barber	67	160455
4. Mr. Euphile's Barber	171	152980
5. Mr. Euphile's Barber	128	173992
6. Mr. Euphile's Barber	67	112729
7. Mr. Euphile's Barber	520	162948
8. Mr. Euphile's Barber	67	10118
9. Mr. Euphile's Barber	108	2060
10. Mr. Euphile's Barber	824	2664
11. Mr. Euphile's Barber	803	860
12. Mr. Euphile's Barber	243	6595
13. Mr. Euphile's Barber	232	1294911
14. Mr. Euphile's Barber	110	610910
15. Mr. Euphile's Barber	102	102907
16. Mr. Euphile's Barber	102	102919
17. Mr. Euphile's Barber	124	140971
18. Mr. Euphile's Barber	339	10797
19. Mr. Euphile's Barber	119	1090
20. Mr. Euphile's Barber	101	101146
21. Mr. Euphile's Barber	229	1044118
22. Mr. Euphile's Barber	123	66673
23. Mr. Euphile's Barber	128	10693
24. Mr. Euphile's Barber	119	25911
25. Mr. Euphile's Barber	420	170295

(a) The PageXML data of our ground truth plotted on part of the original image.



(b) The PageXML data of the HTR system plotted on part of the original image.

Figure 10: The PageXML data of the ground truth and HTR system plotted on the original image. The tabular structure of the ground truth values can be seen where this is not the case in the HTR output.

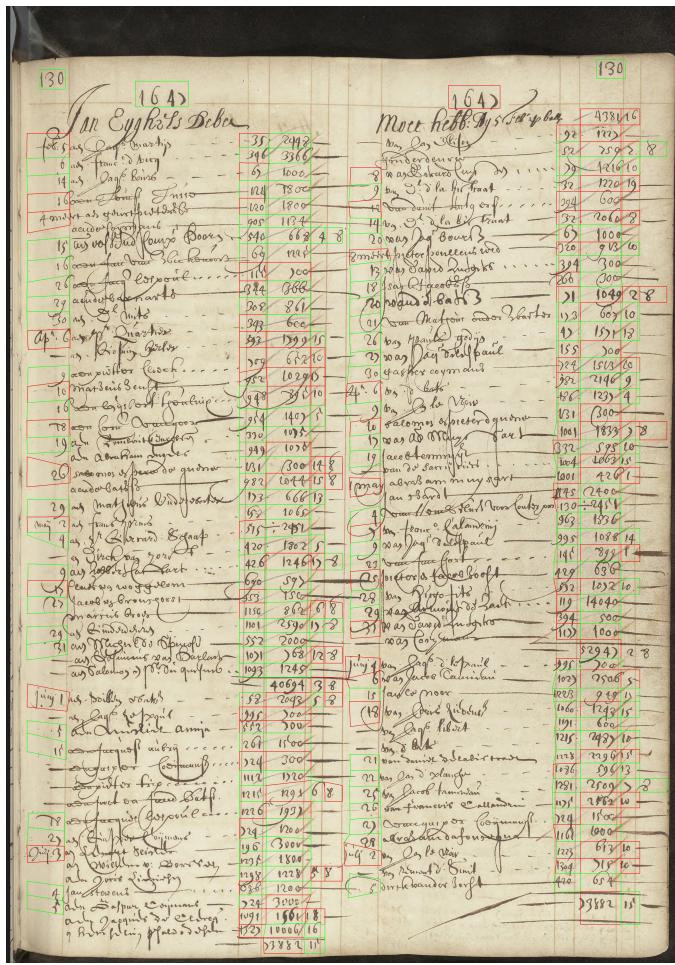


Figure 11: An example page showing the words correctly matched by the HTR system in green and the words wrongly matched in red. This can be helpful in visually inspecting the page to see whether there are common characteristics on parts of the page that could cause mistakes.

607 Appendix B DETAILED METHODOLOGY

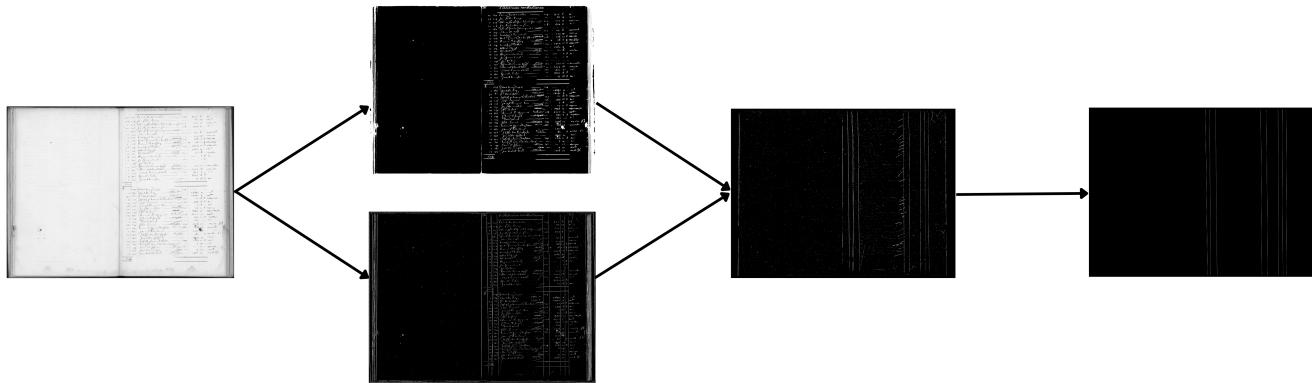


Figure 12: The process of the extraction of column lines. We start with the blurred grayscale image. From there we create two threshold images, one containing the faint column lines and one that does not contain them. We subtract the images from each other, remove noise and use Hough line detector to find the columns lines.



Figure 13: The process of extracting the rows from the ledger page. We use the HTR system to determine bounding boxes of words from the original page. We then determine the row bounding boxes based on the y-coordinates of the word bounding boxes.

608 **Appendix C COMPLEMENTARY RESULT TABLES**

Table 6: Table Cell Performance Metrics at IoU Threshold 0.5.

Subdir	IoU Score	Precision	Recall	F1
1045	0.638	0.674	0.674	0.674
1390	0.641	0.303	0.468	0.363
62	0.627	0.437	0.487	0.460
149	0.608	0.310	0.383	0.342
1099	0.621	0.532	0.499	0.515
99	0.638	0.525	0.504	0.514
109	0.645	0.520	0.552	0.534
119	0.617	0.422	0.437	0.429
139	0.626	0.435	0.475	0.453
59	0.664	0.555	0.541	0.548
89	0.632	0.466	0.519	0.490
69	0.626	0.391	0.456	0.420
55	0.631	0.287	0.319	0.302
79	0.614	0.389	0.439	0.412
1016	0.661	0.637	0.635	0.636
199	0.594	0.243	0.339	0.278
129	0.617	0.342	0.404	0.369
Mean	0.627	0.416	0.459	0.434

Table 7: Table Cell Performance Metrics at IoU Threshold 0.25.

Subdir	IoU Score	Precision	Recall	F1
1045	0.546	0.990	0.990	0.990
1390	0.498	0.513	0.785	0.613
62	0.471	0.854	0.959	0.902
149	0.440	0.710	0.879	0.784
1099	0.473	0.939	0.881	0.909
99	0.504	0.916	0.878	0.896
109	0.508	0.848	0.894	0.867
119	0.460	0.869	0.901	0.885
139	0.474	0.819	0.894	0.853
59	0.525	0.910	0.888	0.899
89	0.475	0.862	0.961	0.907
69	0.464	0.821	0.963	0.884
55	0.451	0.726	0.811	0.765
79	0.452	0.829	0.929	0.875
1016	0.553	0.835	0.831	0.832
199	0.432	0.639	0.894	0.731
129	0.453	0.738	0.878	0.801
Mean	0.473	0.804	0.898	0.844

Table 8: Table Cell Performance Metrics at IoU Threshold >0

Subdir	IoU Score	Precision	Recall	F1
1045	0.196	1.000	1.000	1.000
1390	0.152	0.603	0.952	0.729
62	0.190	0.876	0.984	0.925
149	0.169	0.762	0.943	0.841
1099	0.206	0.945	0.886	0.915
99	0.177	0.960	0.920	0.939
109	0.188	0.885	0.933	0.905
119	0.178	0.906	0.939	0.922
139	0.185	0.860	0.938	0.896
59	0.192	0.946	0.923	0.934
89	0.186	0.873	0.974	0.919
69	0.195	0.839	0.984	0.903
55	0.172	0.820	0.911	0.862
79	0.182	0.863	0.967	0.911
1016	0.232	0.845	0.840	0.841
199	0.169	0.690	0.962	0.788
129	0.164	0.802	0.952	0.869
Mean	0.182	0.845	0.946	0.887